

Interpretable Features for Explaining Machine Learning Models

Mateusz Staniak, MI² Data Lab @ Warsaw University of Technology

Joint work with Przemysław Biecek

Leuven, 15 VII 2019

The need for interpretability

Before

5 years of web logs + ML
=

proved to be a more useful and timely indicator [of flu] than government statistics with their natural reporting lags

- Viktor Mayer-Schönberger and Kenneth Cukier ,
Big Data: A Revolution That Will Transform How We Live, Work and Think

After

WHAT WE CAN LEARN FROM
THE EPIC FAILURE OF GOOGLE
FLU TRENDS




wired.com/2015/10/can-learn-epic-failure-google-flu-trends/

Machine Learning models are vulnerable to:

- Biased training (and other data quality issues),
- Concept drift,
- Unmeasurable objectives (Fairness, Lawfulness).



Females vs Male frequencies



Training data vs validation data

Amazon scraps secret AI recruiting tool that 'didn't like women'

- Amazon ended job recruiting service that was reportedly biased against women
- It was created by Amazon's Edinburgh team in 2014 to automatically sort CVs
- The AI taught itself to downgrade resumes that included words like 'women's'

Types of Explanations

Intrinsic vs post-hoc

- Intrinsic explanations are based on specific algorithm design (model form, training explanations and model jointly).
- Post-hoc explanations are concerned with an already trained model (model is never re-fitted, only predictions are used).

```
lm(formula = m2.price ~ ., data = apartments)
```


Residuals:

Min	1Q	Median	3Q	Max
-247.5	-202.8	-172.8	381.4	469.0

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)	
(Intercept)	5020.1391	682.8721	7.352	4.11e-13	***
construction.year	-0.2290	0.3483	-0.657	0.5110	
surface	-10.2378	0.5778	-17.720	< 2e-16	***
floor	-99.4820	3.0874	-32.222	< 2e-16	***
no.rooms	-37.7299	15.8440	-2.381	0.0174	*
districtBielany	17.2144	40.4502	0.426	0.6705	
districtMokotow	918.3802	39.4386	23.286	< 2e-16	***
districtOchota	926.2540	40.5279	22.855	< 2e-16	***
districtPraga	-37.1047	40.8930	-0.907	0.3644	
districtSrodmiemie	2080.6110	40.0149	51.996	< 2e-16	***
districtUrsus	29.9419	39.7249	0.754	0.4512	
districtUrsynow	-18.8651	39.7565	-0.475	0.6352	
districtWola	-16.8912	39.6283	-0.426	0.6700	
districtZoliborz	889.9735	40.4099	22.024	< 2e-16	***

signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1





Model-agnostic Approach

	Definition	Example	Comments
Model-agnostic explanations	Do not use knowledge about the specific algorithm	Permutation-based variable importance	Do not require model re-fitting
Model-specific explanations	Assume that a specific algorithm was used to fit the model	Average minimum depth in a random forest	Can be more accurate

Local vs Global Explanations

- Local explanations are concerned with a single observation and its prediction.
- Global explanations are concerned with the model as a whole.

Global explanations are often aggregation of local explanations (e.g. mean).




Local Explanations

Approaches to Local Explanations

- What-If analysis (marginal response of the model when changing a single variable for a single observation):
 - Ceteris Paribus profiles (Individual Conditional Expectation).
- Local surrogate models (fitting an interpretable model locally – original paper: LIME, 2016):
 - LIME and its modifications (aLIME, k-LIME, localSurrogate),
 - **LIVE and localModel (versions of LIME developed at MI2 Data Lab).**
- Example-based explanations
 - Contrastive explanations
 - Prototypes and criticism
- Prediction decomposition (attributing additive scores to features).
 - EXPLAIN,
 - Shapley Values,
 - **Break Down and iBreakDown (methods connected to Shapley Values developed at MI2 Data Lab)**

Local Interpretable Model-agnostic Explanations





[1]

M. T. Ribeiro, S. Singh, C. Guestrin, „«Why Should I Trust You?»: Explaining the Predictions of Any Classifier”, *Proceedings of the 22Nd ACM SIGKDD International Conference on Knowledge Discovery and Data Mining*, 2016

- Optimization problem:

$$\mathcal{L}(f, g, \pi_x) = \sum_{z, z' \in \mathcal{Z}} \pi_x(z) (f(z) - g(z'))^2$$

- f is the explained model,
- g is the explanation model,
- z is a interpretable representation of x ,
- π is a distance measure (a kernel).



LIME explanations

Multiclass case

Prediction probabilities

atheism	0.50
christian	0.43
religion.misc	0.05
mid-east	0.02
Other	0.00

NOT atheism

Caused	0.26
Rice	0.15
Genocide	0.13
certainty	0.09
scri	0.09
owl-net	0.08

atheism

NOT christian

Caused	0.18
fsu	0.09
Genocide	0.09
scri	0.09
Semitic	0.08
Luthers	0.08

christian

Tabular data

Prediction probabilities

edible	0.00
poisonous	1.00

edible

odor=foul	0.26
gill-size=broad	0.13
stalk-surface-above-ring=silky	0.11
spore-print-color=chocolate	0.08
stalk-surface-below-ring=silky	0.06

poisonous

Feature

Feature	Value
odor=foul	True
gill-size=broad	True
stalk-surface-above-ring=silky	True
spore-print-color=chocolate	True
stalk-surface-below-ring=silky	True



<https://github.com/marcotcr/lime>

1)

X_1	...	X_p
M	...	0.11
F	...	-0.25
...
U	...	0.887

$f(x)$

y
0.87
0.14
...
0.54

Z_1	...	Z_p
$X_1 = M$...	$X_p < 0.2$
$X_1 = F, U$...	$X_p < 0.2$
...
$X_1 = F, U$...	$X_p > 0.2$

$X_1 = F, U$...	$X_p < 0.2$
--------------	-----	-------------

2)

$X_1 = F, U$...	$X_p < 0.2$
--------------	-----	-------------

$X_1 = M$...	$X_p < 0.2$
$X_1 = F, U$...	$X_p > 0.2$
$X_1 = M$...	$X_p > 0.2$
...
$X_1 = F, U$...	$X_p < 0.2$

Z_1	...	Z_p
$X_1 = M$...	$X_p < 0.2$
$X_1 = F, U$...	$X_p > 0.2$
$X_1 = M$...	$X_p > 0.2$
...
$X_1 = F, U$...	$X_p < 0.2$

3)

Z_1	...	Z_p
$X_1 = M$...	$X_p < 0.2$
$X_1 = F, U$...	$X_p > 0.2$
$X_1 = M$...	$X_p > 0.2$
...
$X_1 = F, U$...	$X_p < 0.2$

X_1	...	X_p
M	...	0.111
F	...	1.27
M	...	0.887
...
F	...	-0.2


4)

f(x)

f(x)
0.93
0.77
0.122
...
0.64




5)

z_1	...	z_p	$f(x)$
$X_1 = M$...	$X_p < 0.2$	0.93
$X_1 = F, U$...	$X_p > 0.2$	0.77
$X_1 = M$...	$X_p > 0.2$	0.122
...
$X_1 = F, U$...	$X_p < 0.2$	0.64



$g(z)$
0.90
0.81
0.07
...
0.641

6)




Some remarks on LIME for tabular data

- Most of the work so far focused on step **2)** – the sampling:
 - Laugel et al. 2018: the neighbourhood must include the decision boundary,
 - Adhikari et al. 2018: the neighbourhood must include enough data points from both classes,
 - Tan et al. 2019: sampling introduces significant uncertainty.
- For image or text step **3)** is trivial, but for tabular data and non-trivial interpretable input spaces, the inverse transformation is a problem.
- For tabular data, step **1)** is important, but often features are not transformed. It is not clear, what should be considered an interpretable feature.

Existing Approaches for Tabular Data

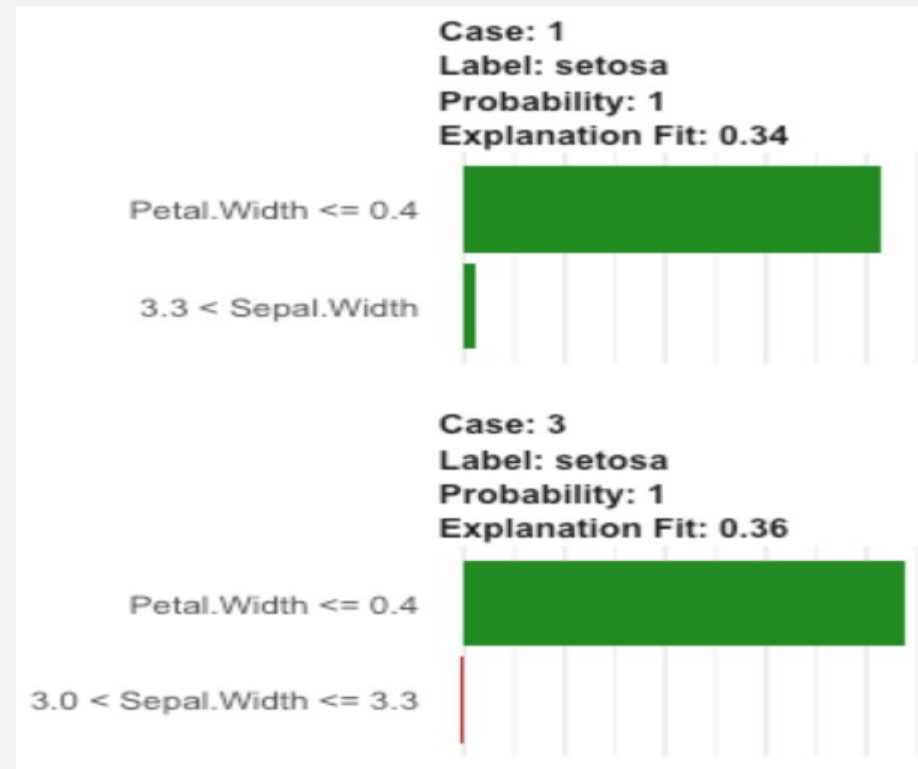
Discretized features

- lime library (Python)
 - Marco Tulio Ribeiro



<https://github.com/marcotcr/lime>

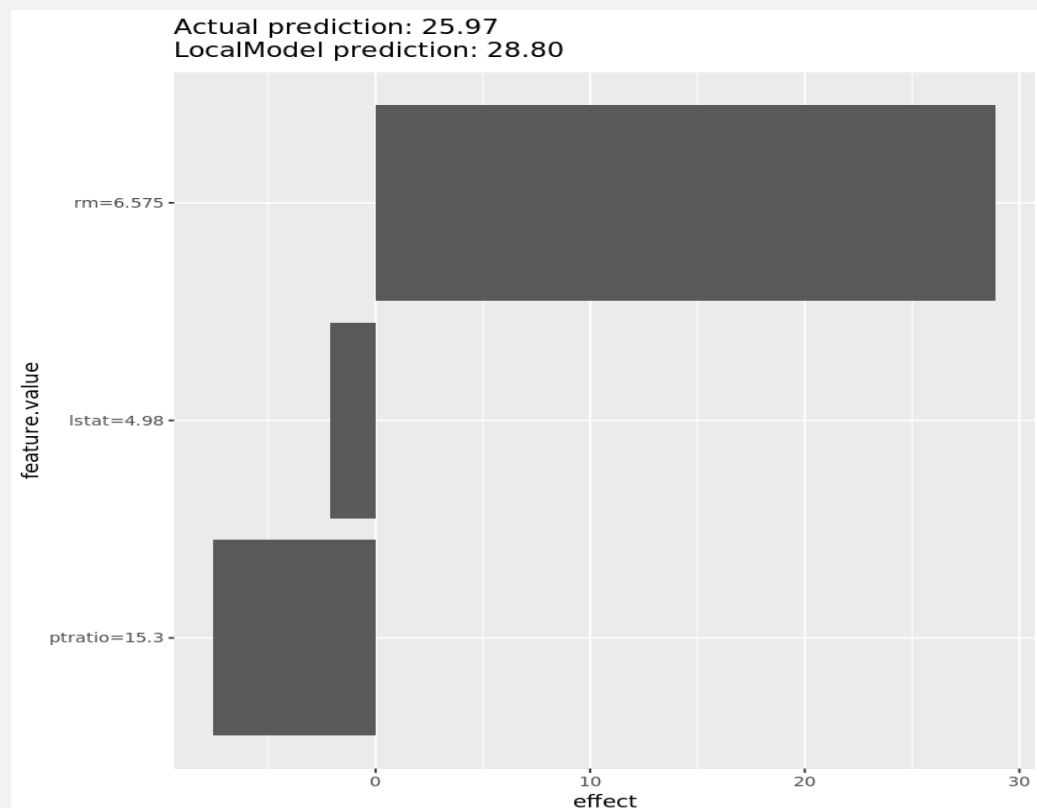
- lime package (R)
 - Thomas Lin Pedersen



<https://github.com/thomasp85/lime>

Continuous features

- iml package (R) – Christoph Molnar (JOSS, 2018)



<https://github.com/christophM/iml>

- live package (R) – Mateusz Staniak (R Journal, 2018)


Variable	N	Estimate	p
fixed_acidity	2000	0.14 (0.14, 0.15)	<0.001
volatile_acidity	2000	-1.43 (-1.46, -1.40)	<0.001
citric_acid	2000	-0.66 (-0.69, -0.63)	<0.001
residual_sugar	2000	0.00 (-0.00, 0.01)	0.9
chlorides	2000	-2.57 (-2.71, -2.43)	<0.001
free_sulfur_dioxide	2000	0.00 (0.00, 0.00)	<0.001
total_sulfur_dioxide	2000	0.00 (0.00, 0.00)	<0.001
density	2000	-25.69 (-28.09, -23.30)	<0.001
pH	2000	-0.82 (-0.85, -0.79)	<0.001
sulphates	2000	2.56 (2.53, 2.60)	<0.001
alcohol	2000	0.20 (0.19, 0.20)	<0.001
(Intercept)		30.32 (27.93, 32.71)	<0.001

<https://github.com/MI2DataLab/live>


A New Approach:


Use our knowledge about the model
behaviour


Partial Dependence Plots




Ceteris Paribus Profiles









Summary

- Explanations of individual predictions rely on good interpretable features.
- For tabular data, the notion of an *interpretable feature* is not clear.
- We propose a method of creating interpretable features based on conditional behaviour of the model. It is implemented in R package **localModel**.

References

- Biecek, P., 2018. DALEX: explainers for complex predictive models, Journal of Machine Learning Research 19(84):1–5, 2018.
- Goodman, B., Flaxman, S., 2017. European Union regulations on algorithmic decision-making and a “right to explanation.” AI Magazine 38, 50. <https://doi.org/10.1609/aimag.v38i3.2741>
- Goldstein, A., Kapelner, A., Bleich, J., Pitkin, E., 2013. Peeking Inside the Black Box: Visualizing Statistical Learning with Plots of Individual Conditional Expectation. arXiv:1309.6392 [stat].
- Lundberg, S.M., Lee, S.-I., 2017. A Unified Approach to Interpreting Model Predictions, Advances in Neural Information Processing Systems 30. Curran Associates, Inc., pp. 4765–4774.

References

- Molnar, Christoph. "Interpretable machine learning. A Guide for Making Black Box Models Explainable", 2019. <https://christophm.github.io/interpretable-ml-book/>.
- Ribeiro, M.T., Singh, S., Guestrin, C., 2016. "Why Should I Trust You?": Explaining the Predictions of Any Classifier.
- Robnik-Šikonja, M., Bohanec, M., 2018. Perturbation-Based Explanations of Prediction Models, in: Zhou, J., Chen, F. (Eds.), Human and Machine Learning: Visible, Explainable, Trustworthy and Transparent. Springer International Publishing, Cham, pp. 159–175. https://doi.org/10.1007/978-3-319-90403-0_9
- Staniak, M., Biecek, P., 2018. Explanations of model predictions with live and breakDown packages. Mateusz Staniak and Przemysław Biecek , The R Journal (2018) 10:2, pages 395-409.
- Miller, Tim. "Explanation in artificial intelligence: Insights from the social sciences." arXiv Preprint arXiv:1706.07269. (2017)

References

- T. Laugel, X. Renard, M.-J. Lesot, C. Marsala, i M. Detyniecki, „Defining Locality for Surrogates in Post-hoc Interpretability”, *arXiv:1806.07498 [cs, stat]*, cze. 2018
- A. Adhikari, D. M. J. Tax, R. Satta, i M. Fath, „Example and Feature importance-based Explanations for Black-box Machine Learning Models”, *arXiv:1812.09044 [cs]*, grudz. 2018.
- D. Alvarez-Melis i T. S. Jaakkola, „On the Robustness of Interpretability Methods”, *arXiv:1806.08049 [cs, stat]*, cze. 2018.
- H. Fen, Tan, K. Song, M. Udell, Y. Sun, i Y. Zhang, „Why should you trust my interpretation? Understanding uncertainty in LIME predictions”, *arXiv:1904.12991 [cs, stat]*, kwi. 2019.
- Molnar, C., Bischl, B., Casalicchio, G., 2018. iml: An R package for Interpretable Machine Learning. JOSS 3, 786. <https://doi.org/10.21105/joss.00786>